

(a) above-onset realisation

PENDING

fringe-marginalised $p(L|d)$
concentrated on the correct fringe

*SAMPLER g3: NOT RUN
(no chain banked; cronus
OOM, ~7 GPU-hr/chain)*

(b) census loudness

PENDING

same machinery — honestly
multimodal/diffuse
(concentration tracks evidence)

*SAMPLER g3: NOT RUN
(stated prediction only,
SAMPLER_dev.md §4)*

(c) σ_{MLE} VS σ_{sampler}

PENDING

linear-parameter covariance
underestimate, per parameter

*no such comparison in
the SAMPLER record;
requires the unrun chains*